ChatGPT

Dynamic Gamma, Vanna, and Charm Fields in a Time–Price Landscape

Introduction

In option markets, the sensitivity of option dealers' positions to underlying price moves, volatility shifts, and time decay can be visualized as dynamic **fields** overlaying the stock's time–price plane. In this framework, **time** is treated as the x-axis and **stock price** as the y-axis, while three Greek-derived fields – **gamma (Φ_Γ)**, **vanna (Φ_{Vanna})**, and **charm (Φ_{Charm})** – represent "pressure" exerted on the market at each point in time and price. These fields quantify how dealer hedging flows evolve, potentially influencing past and present price stability and even conditioning future movements. In this report, we explore the theoretical foundations of these Greeks as partial derivatives in the Black–Scholes model, then examine empirical and institutional perspectives on how gamma, vanna, and charm flows create feedback loops in real markets. We discuss how these fields interact with liquidity, implied volatility, and different market regimes, and how modern trading systems implement such models (e.g. constructing an "implied order book" and dynamic delta hedging engines). Finally, we consider known limitations and pitfalls of treating option Greeks as continuous fields. The goal is a deep technical and conceptual analysis, drawing on both academic research and practitioner insights.

Greeks as Dynamic Fields: Theoretical Foundations

Option Greeks are partial derivatives of option value, and they naturally lend themselves to interpretation as *rates of change* along the dimensions of the time–price landscape. Gamma (Γ) measures the rate of change of delta with respect to the underlying's price; vanna is a second-order cross-sensitivity measuring how delta changes with volatility (often considered ∂Δ/∂σ or equivalently ∂Vega/∂S); and charm (sometimes called delta decay) measures the rate of change of delta with respect to time [1] [2]. In Black–Scholes, closed-form expressions for these Greeks can be derived. For example, the formula for gamma of a call or put is:

$$\Gamma = \frac{e^{-q\tau} n(d_1)}{S\,\sigma\sqrt{\tau}}\,,$$

where $n(d_1)$ is the standard normal density, q is the dividend yield, σ the implied volatility, τ time to expiration, and S the spot price [3]. A key property is that gamma is always positive for calls and puts (long options), indicating convexity in the price dimension. **Vanna** can be defined as the derivative of an option's delta with respect to volatility (or equivalently, ∂Option/∂S∂σ). Under Black–Scholes assumptions, one expression is:

$$\mathrm{Vanna} = -\,e^{-q\tau}\,n(d_1)\,\frac{d_2}{\sigma}\,,$$

which captures how delta (or option value) will shift if implied volatility changes [4]. **Charm**, or delta decay, is the partial derivative of delta with respect to the passage of time (with S and σ held constant). Its formula is a bit more complex, but essentially charm for an option is negative for calls and puts – meaning delta tends to decrease in magnitude as time passes, all else equal [5]. Charm is often expressed per day or per minute; for instance, converting the annualized formula to per-minute units ensures consistent dimensions in models [6] [7].

Crucially, these Greeks can be *aggregated across all option positions held by dealers* to produce field values at a given time. Imagine a surface over the (time, price) plane where at each point (t, S) we plot Φ_Γ(t,S): the total dollar gamma exposure of dealers (signed by position) at that time if the underlying were at price S. Similarly, Φ_{Vanna}(t,S) would represent the total sensitivity of dealer delta to a small volatility change, and Φ_{Charm}(t,S) the rate at which dealers' delta exposure is changing with time. These "pressure fields" are dynamic: as time progresses (moving along x-axis), options decay (altering charm), and as the underlying price moves along the y-axis, gamma and vanna exposures shift (because options go more in- or out-of-the-money).

To compute such fields in practice, one must sum the contributions of thousands of options. Each option i contributes a dollar gamma $G_{\$,i} = \Gamma_i \, S^2 \, OI_i \, M$ (where OI_i is the number of contracts and M the contract multiplier) [8]. Summing $G_\$$ for all options (with a sign depending on whether dealers are net long or short each option) yields the aggregate Φ_Γ in dollars per 1-point underlying move [9] [10]. Likewise, summing each option's "vanna dollars" $V_\$ = \mathrm{Vanna}_i \, S \, OI_i \, M$ (often weighted by the correlation between spot and vol moves) gives Φ_{Vanna} [10]. Charm can be summed in units of $/minute (or $/day) as $C_\$ = \mathrm{Charm}_i \, S \, OI_i \, M$ [11] [12]. The result is a set of fields: $\Phi_\Gamma(t, S)$ (in $ per ΔS), $\Phi_{Vanna}(t, S)$ (in $ per Δσ), and $\Phi_{Charm}(t, S)$ (in $ per minute) that vary over time and price.

Mathematically, these fields represent partial derivatives in the *landscape* of the option market's risk profile. Gamma is like a curvature along the price axis (second derivative w.rt price), vanna mixes the price and volatility directions, and charm is a gradient along the time axis. By overlaying them on a time–price plot, we can visualize pockets of high "pressure" – for example, regions where dealers have concentrated gamma exposure at certain strikes and expirations, or where imminent expiries are causing large charm effects. These regions indicate where and when dealer hedging flows could be significant forces in the market.

Gamma Field and its Market Impact

Gamma is the most well-known Greek in terms of impact on underlying price dynamics. It measures how an option's delta changes for small moves in the underlying. A dealer who is *long gamma* (typically from being long options or having sold options to customers who are net short) will have a delta that increases when the market rises and decreases when the market falls. To remain hedged (delta-neutral), the dealer must *sell* into rising markets and *buy* into falling markets, acting as a counterbalancing force. This behavior **supplies liquidity and dampens volatility**: the dealer's trades oppose the price direction, tending to mean-revert the market [13]. Conversely, a dealer who is *short gamma* (having sold options to customers who are net long options) finds that as the market rises, their delta *falls* (because short option positions lose delta), forcing them to **buy** into an advance; as the market falls, their delta rises (becoming more negative for short calls or more positive for short puts), forcing them to **sell** into the decline [14]. In a short gamma regime, dealer hedging **removes liquidity and amplifies volatility**, as their trades reinforce the market's move. These dynamics are illustrated in **Figure 1**, which shows how a dealer's net gamma exposure dictates hedge flows: positive gamma positions result in stabilizing sell-on-strength and buy-on-weakness flows, while negative gamma positions result in de-stabilizing buy-on-strength and sell-on-weakness flows [13].

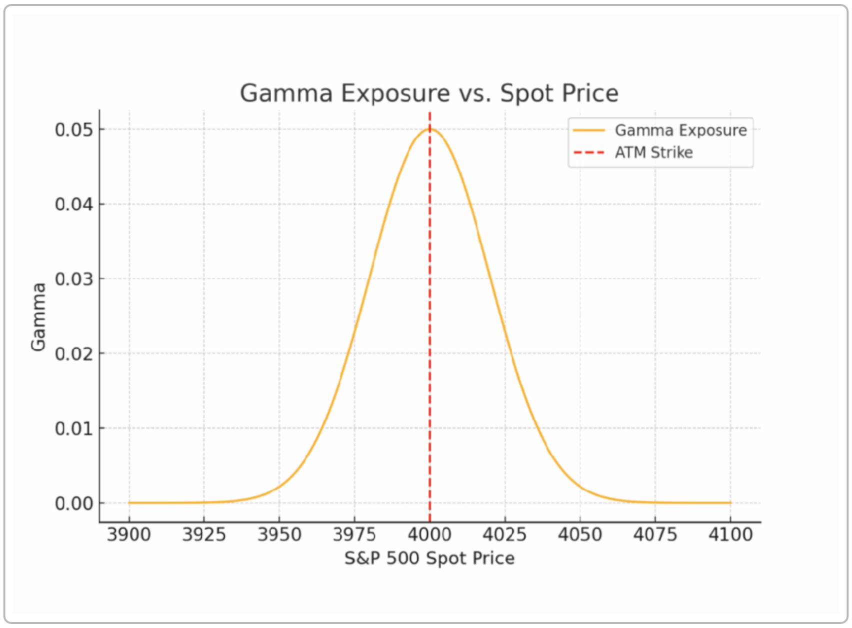

Figure 1: Gamma Exposure vs. Spot Price – illustration of a dealer's aggregate gamma exposure concentrated around the at-the-money level. When the market is near the high-gamma strike (red dashed line), small price moves create significant delta changes (peak gamma), meaning large hedging flows. Away from that region, gamma (and thus required hedge size per point) tapers off, indicating lower pressure. [13]

Because index options markets often have strikes with massive open interest, one can map out gamma exposure versus price to identify "Gamma peaks" or **pinning points**. Near option expiry, gamma for at-the-money strikes tends to explode (since $\Gamma \sim 1/(\sigma\sqrt{\tau})$, with $\tau \to 0$), so if the underlying trades near a large strike, dealers' gamma position can exert a magnetic pull on the price. Market observers refer to this as **pin risk** or "gamma pinning" – the price gravitating toward a strike where gamma is highest, as dealers aggressively hedge even tiny moves. If the aggregate gamma position is long (dealers long gamma), the price may get stuck (low volatility) near that strike; if dealers are short gamma at the strike, the price might violently sling away once it starts moving beyond the range (since hedging then accelerates the move) [15]. This behavior is often seen around large quarterly options expirations, where the *gamma field* is concentrated at a few strikes and time (expiry) is imminent.

Empirical evidence strongly supports gamma's role in modulating market volatility and momentum. A study by Societe Generale noted that **all the large daily S&P 500 moves (>1.5%) in one sample period occurred on days when the previous day's aggregate gamma exposure was *negative*** [16]. In contrast, when dealers are long gamma (aggregate gamma positive), intraday price action tends to be muted as their hedging squelches directional movement [17]. Academic research confirms this effect: simulations by Buis et al. (2023) found that increasing the net gamma position of dynamic hedgers **reduces volatility and enhances market stability**, whereas a negative net gamma position **increases volatility and makes the market more prone to failures** [18]. Moreover, price discovery can suffer in extreme gamma regimes – when hedgers dominate, prices may reflect hedging flows more than fundamentals [19].

Recent empirical work on the explosion of **0DTE (zero days to expiry)** options (which are effectively all gamma and charm, given their ultra-short life) shows that market makers' net gamma intraday is usually positive, and this *positive gamma correlates with lower subsequent volatility and intraday mean-reverting price action*. Conversely, on occasions of net short gamma, the market exhibits momentum (trend continuation) and higher volatility [20]. In other words, even on a very short timescale, gamma's influence is consistent with classic theory: positive gamma inventory leads dealers to counteract price swings (promoting reversals), while negative gamma leads them to chase price moves (promoting momentum). This underscores gamma's importance as a **field of market "pressure"** – when the gamma field is positive (elevated in the region of current price), it acts like a stabilizing pressure blanket on the market; when the gamma field is negative, it's like a slippery slope that can accelerate movements.

It's useful to think of the gamma field as forming an **implied liquidity profile** for the market. If we sum up how many dollars of the underlying dealers would need to buy or sell for each small increment of price (i.e. the gamma exposure at each price level), it resembles a limit order book. Areas with large positive gamma indicate *abundant liquidity* (dealers will provide lots of opposing flow if price moves there), whereas areas with large negative gamma indicate a potential *liquidity vacuum* or even self-reinforcing flow. In fact, researchers have described **options as an "enigmatic order type" that reveals the true liquidity landscape**: by mapping where delta hedges must occur, one can **build a uniquely information-rich *"implied order book"*** showing supply (liquidity) and demand holes in the market [21]. Unlike the visible order book (often thin or deceptive in modern markets), this implied order book from options is more transparent because it's derived from committed positions and the *necessity* of hedging them [22] [23]. This concept was popularized by practitioners (e.g. SqueezeMetrics) who note that the *change* in option delta – driven by underlying moves (gamma), time decay (charm), or volatility shifts (vanna) – compels dealers to trade in a predictable way, much like resting orders that materialize as price moves [23]. In practice, aggregating the gamma field across all strikes gives a profile of how many shares dealers would buy or sell *per 1-point move* at various index levels [24] [25]. A positive aggregate gamma means dealers as a whole would *provide* shares on both sides (sell into rises, buy into dips), stabilizing the market [24] [26]; a negative aggregate gamma means dealers will *take* liquidity (buy into rises, sell into dips), potentially destabilizing the market. The gamma field thus connects directly to observable outcomes: for instance, a strongly positive gamma field often corresponds to periods of realized volatility being much lower than implied volatility, as dealer hedging dampens fluctuations [17]. On the other hand, when the gamma field turns negative (often during market sell-offs when put positions dominate and dealers are short gamma), volatility can far exceed expectations and even contribute to crash-like dynamics.

Vanna Field and Volatility Feedback

While gamma relates to directional moves, **vanna** captures the interplay between volatility and delta. Vanna is the sensitivity of an option's delta to changes in implied volatility (and often practically measured as the sensitivity of option value to a simultaneous change in spot and vol). In plainer terms, if implied volatility (IV) shifts up or down, how do dealers' delta exposures change? This is crucial because in stress scenarios, underlying prices and volatility often move inversely (stocks down, volatility up – the classic "volatility feedback"). The **vanna field** $\Phi_{Vanna}(t, S)$ maps how much delta hedging flow would be triggered by a volatility change at a given time/price.

Practitioners sometimes call vanna the "**volatility feedback engine**" of the market [27]. If dealers have significant vanna exposure, a change in implied volatility will force them to adjust their hedges, which can in turn affect the underlying price, which feeds back into volatility. This can set off a **feedback loop**. For example, consider a common positioning: a customer hedges downside by **buying puts and selling calls** (a risk-reversal), and a dealer takes the other side (dealer is short puts, long calls). Initially, the dealer might delta-hedge so that the position is neutral. Now suppose the index rallies moderately. The short put becomes far out-of-the-money (its delta drops), and the long call gets closer to the money (its delta rises). The dealer's net delta becomes long (due to the call), so they **sell futures** to rebalance. Additionally, the rally likely causes implied vol to **fall** (volatility typically declines in rallies). When vol falls, both the put and call deltas decrease (options become less sensitive when IV is lower), which means the dealer's previous hedge – the futures they sold – is now too much. To correct this **over-hedge**, they **buy back** futures [28] [29]. In a falling market scenario, the opposite happens: underlying drops, dealer buys to hedge, vol spikes, making deltas larger, and suddenly the dealer is *under*-hedged (needing to sell more). In both cases, there's a **back-and-forth** hedging induced by the coupled changes in spot and vol. This *vanna-driven loop* can reinforce the market's trend: in a rally, dealer hedging (sell then buy) can actually add **extra buying on the second leg**, helping propel the rally further and suppressing volatility more; in a decline, the loop can add extra selling and exaggerate the vol spike. As one analyst put it, the hedging drives price, which drives volatility, which reshapes hedge needs, *and so on*, forming a self-reinforcing cycle [29]. In high-volume or large-position scenarios, such **vanna feedback loops can lead to runaway price moves disconnected from fundamentals** [29].

It's important to note that **vanna's effect depends on the *moneyness* of positions and the direction of volatility change**. Unlike gamma (which for a given option always pushes in one liquidity direction or the other depending on the sign of the position), vanna can flip the sign of its impact based on context. For instance, consider a single long put position held by a dealer (meaning the customer sold the put to the dealer). If that put is out-of-the-money (OTM), an increase in implied vol will *increase* the put's delta (making it less OTM in a probabilistic sense) – in the example from SqueezeMetrics, a 2900 strike SPX put with 1 month to expiry sees its delta rise from 0.27 to 0.30 when IV jumps from 20% to 25%, forcing the dealer to buy more index to hedge [30] [31]. But if the put were in-the-money (ITM), raising IV actually *lowers* its delta (because higher vol increases the chance it might end out-of-the-money), meaning the dealer would **sell** some of their index hedge when vol rises [32]. Thus, the same vol increase can cause dealer *buying* or *selling* depending on the moneyness of the option (OTM vs ITM) – *"for Vanna, moneyness matters."* [33] This complexity means the vanna field is often patchy and harder to predict; one must consider the overall skew of positions. Typically, in index markets, customers tend to be long OTM puts and short OTM calls (for protection and yield strategies) [34]. Both of these position types have the effect that **when IV rises (usually as the market falls and liquidity worsens), dealers must *sell* underlying** (since dealers are short puts and long calls in those scenarios, both of which lead to selling when vol jumps) [35]. This is an inherently **unstable feedback**: a market decline begets a vol spike, which triggers dealer selling, which further pressures the market, etc. In the colorful phrasing of SqueezeMetrics, **vanna is gamma's "evil twin."** Gamma, as we saw, generally supplies liquidity and keeps volatility in check *until it doesn't* – and what often causes it to fail is vanna-driven flow overwhelming gamma's stabilizing influence [34] [35].

Empirical data backs up the critical role of vanna in crises. Historical reconstructions of S&P 500 **vanna exposure (VEX)** show that it can reach large negative values in major sell-offs. In late 2008, for example, the aggregate VEX was about **–\$400 million per 1% change in the S&P 500** (meaning for each 1% increase in implied vol or corresponding large drop in the index, dealers would need to *sell \$400mm* of S&P 500 to re-hedge) [36]. A similarly extreme negative vanna was seen during the March 2020 COVID crash [36]. It is no coincidence that these periods – with **deeply negative vanna exposure** – were among the most volatile and illiquid in recent memory. When the vanna field turns negative and large in magnitude, it signifies that the "derivative tail" can wag the market dog: dealer hedging demand from volatility shifts is so strong that it fuels sustained volatility in the underlying [37] [38]. In effect, the volatility market and spot market become a two-way feedback machine, with dealers caught in the middle.

In calmer times, however, vanna can work in the market's favor. During periods of declining implied volatility (vol crush) – often when markets grind higher or after a risk event passes – vanna means dealers *reduce* their hedges and thereby add a mild bid to the market. For instance, if customers had previously bought puts (dealer short puts) and implied vol comes down, those puts lose delta, and the dealer can buy back some of their short-futures hedge, contributing to a rally. This is sometimes called a "**vanna rally**." Market commentators often point out that when a volatility event passes and IV drifts lower, dealers' vanna unwinding flows can support equity markets for days or weeks, as long as the underlying doesn't move violently in the opposite direction. Thus, vanna's impact is highly regime-dependent: it's generally a **mild second-order effect** in stable, low-vol markets, but becomes **first-order critical** in fast-moving, high-vol markets when spot-vol correlation is high (e.g. market down, vol up). Modern quantitative models even incorporate a **time-varying correlation, ρ(t)**, between spot and vol to weight vanna's impact – acknowledging that if price and volatility are decoupled, vanna flows might not materialize, but if they are tightly negatively correlated (as in crises), the vanna field's influence on price is amplified [10] [39].

Charm Field and Time-Decay Effects

Charm, also known as **delta decay**, represents the change in an option's delta purely due to the passage of time *t* (with spot and vol unchanged). It reflects how options' probabilities of finishing in or out of the money evolve as the clock ticks down. As expiration nears, in-the-money options' deltas tend to drift toward 1 (for calls) or –1 (for puts), while out-of-the-money options' deltas drift toward 0. This natural progression forces dealers to adjust hedges even if the underlying price doesn't move at all. The **charm field** $\Phi_{Charm}(t, S)$ can be visualized as a steady "flow of delta" over time – effectively a source or sink term on the time axis that causes hedging pressure. Although charm effects are generally smaller than gamma or vanna on any given day, they are continuous and *predictable*, and near major expiries they can have tangible impact on market direction.

If dealers are net short options (short calls or short puts) that are near-the-money, charm will cause their delta exposure to increase as expiry approaches – requiring them to trade in the direction of the underlying movement to remain hedged. Conversely, if dealers are net long options, their delta exposure will shrink over time, requiring opposite-direction hedging. One way to think of it: charm can create a **drift** in the market due to hedging flows, even in the absence of any external news or price trigger. For example, imagine it's the morning of expiration day and a dealer is **short 100 at-the-money put options**, each with a delta of about 50%. At that moment, the dealer is short 50 delta * 100 = 5,000 delta (effectively long 5,000 shares if considering delta from the dealer perspective, since short a put is equivalent to long delta). To hedge this, the dealer sells 5,000 shares of the underlying. Now, as the day progresses toward expiration, one of two things will happen: if the underlying stays around that strike (ATM), the uncertainty will diminish and the put's delta will start moving either towards 1 or 0 by the closing bell – essentially, by expiration it will either be exercised (delta –1 for the short put, which is –100% for the underlying) or expire worthless (delta 0). Let's say the underlying drifts slightly downward, keeping the put ATM to slightly ITM. By late day, each put's delta might have risen from 50% to, say, 80%. Now the dealer's short 100 puts carry 8,000 aggregate delta (short), meaning to stay hedged they must **buy back 3,000 shares** (since they were short 5,000 and now need to be short only 8,000 *puts* delta worth of shares, which is 8,000 short delta offset by short 8,000 shares; initially they had short 5,000 shares). In this scenario, as time ran out, the dealer's need to hedge generated a persistent **buying pressure throughout the day**. More generally, *charm injects a* *consistent bid or offer throughout the day* depending on dealers' position skew [40]. If dealers are short calls or puts, charm tends to make them buy gradually (calls decay -> dealer long calls lose delta, dealer needs to buy back short hedges; puts decay -> dealer short puts lose delta, also needing to buy back hedges). If dealers are long options, charm will have them sell gradually (as their long options lose delta, their long underlying hedges become excessive and are trimmed). In practice, because dealers in index options are often net short puts and short calls (from customers buying puts and selling calls), **charm often creates a mild positive drift** (buying bias) in markets absent other news – a phenomenon sometimes noticeable during "expiration weeks" when, all else equal, indices can grind higher as short-dated options decay and dealers unwind short hedges. One analyst described this as *"Charm, therefore, injects a consistent bid or offer throughout the day – depending on the skew of the book."* [40] If the dealer book is net short puts (downside skew), the charm field is a persistent bid under the market; if net short calls (upside exposure), it could be a persistent sell pressure.

An illustrative example of charm in action is the Monday after a large options expiration. Often a significant portion of dealer delta positions expires on a Friday. By the next trading day, if no new options are traded in equivalent size, the market finds itself "untethered" by those hedging flows – which can sometimes lead to outsized moves (this is one reason why volatility can pick up after big expirations). However, *leading up to* that expiration, charm effects from expiring positions can either pin the market or create a drift. Dealers will be actively readjusting hedges as OTM options decay to zero (releasing hedges) or ATM options start picking direction. Traders watch metrics like "time decay flow" or explicitly compute charm by strike to anticipate how much automatic buying or selling might occur each day due to expiring gamma. One metric sometimes cited is the "charm yield" – e.g. if $X billion of index delta is set to decay over a week (from options losing delta as time passes), then roughly $X billion of underlying would need to be bought (if dealers were short those options) across that week, absent any offsetting moves.

It's worth noting that some practitioners downplay charm compared to gamma and vanna. For example, in the *Implied Order Book* analysis by SqueezeMetrics, charm was deemed "too small an effect to have practical utility" in their liquidity models [41]. Indeed, on longer time scales or for far-dated options, charm's effect on daily trading is minimal. However, for very short-dated options (0DTEs, weekly options) and around key expirations, charm can be locally significant. The charm field is most potent where there is a large aggregate near-the-money position about to expire. In those moments, the time axis pressure (charm) can contribute to sharp moves or intraday trends. For example, on the day of expiry, out-of-the-money options rapidly lose whatever delta they had (if they remain OTM), which can cause a *flurry of hedging unwinds* in the final hour. Market observers often note that the last hour on expiration day – often called "OpEx hour" – can see sudden rallies or drops as remaining delta is shed (sometimes referred to as "charm flows" or "pin dynamics resolving"). In summary, charm is like a gentle breeze pushing the market, usually subtle but directionally predictable, whereas gamma is a sturdier force that resists or accelerates moves, and vanna is a potentially wild gale that can turn a breeze into a storm when volatility shifts.

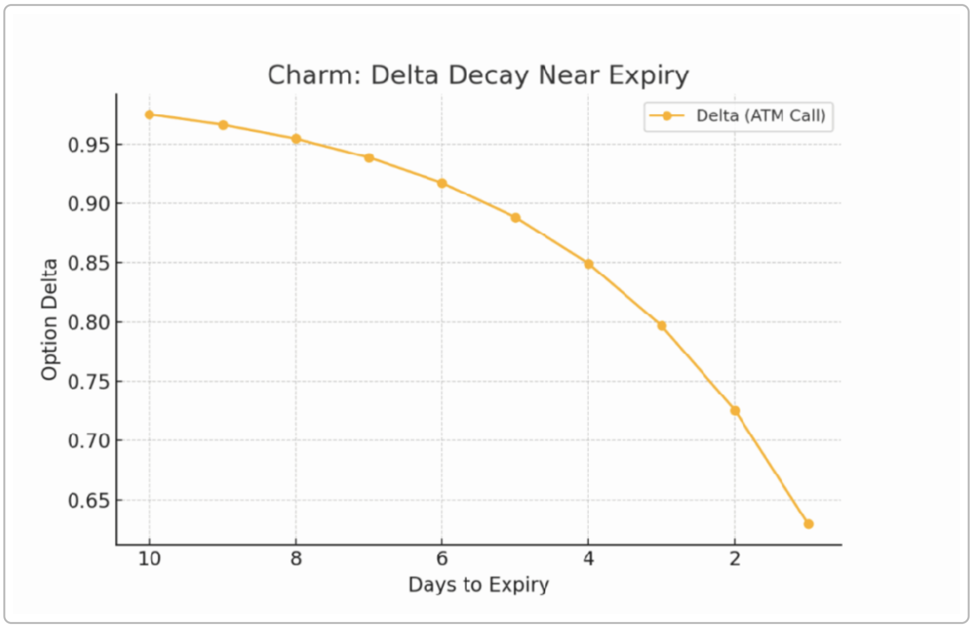

Figure 2: Illustration of Charm (Delta Decay) – As expiration nears, an option's delta will change even if the underlying price remains the same*. This chart shows an example of an (in-the-money) option's delta decaying from ~0.98 to ~0.65 over the last 10 days before expiry. Such delta decay means a dealer long this option would need to reduce their hedge (sell underlying steadily), whereas a dealer short the option (short delta position) would need to buy back underlying as time passes. Charm creates a time-driven hedging flow that acts like a predictable "drift" in the absence of other shocks.* [42] [40]

In quantitative terms, one can integrate the charm field over a time horizon to estimate its cumulative impact. For instance, an *aggregate* Φ_{Charm} of \$5 million per day (dealers' delta is decaying by \$5M worth of index per day, requiring that amount of hedging) over a week would imply about \$25 million of underlying to buy (if dealers short options) or sell (if dealers long) over that week, spread out gradually [43]. Some advanced models define a **"decay wind"** metric to capture this drift: e.g. $D_H = \kappa_t \, w_{CH} \, \Phi_{CH}(t) \times H$, which estimates the price impact (in basis points) over horizon H from charm flows, given an impact factor κ and a weight w_CH [44] [45]. Such metrics treat charm as a sort of gentle wind pushing the price over time, in contrast to gamma/vanna which can cause immediate shocks or amplifications.

Interactions with Liquidity, Volatility, and Market Regimes

The behavior of gamma, vanna, and charm fields does not occur in a vacuum – it is modulated by market liquidity, the state of implied volatility (vol surface), and broader market regimes. A critical concept here is **liquidity sensitivity**: how much does a given amount of delta hedging flow move the underlying price? In a deep, highly liquid market, dealers can adjust hedges with minimal price impact; in a thinner or stressed market, even modest hedging flow can move the market significantly. To incorporate this, sophisticated models use metrics like **Kyle's lambda (λ)** – the price impact per unit of flow [46] [47]. λ(t) can be estimated from market microstructure (e.g. regression of price returns on signed order flow) and typically varies with time of day, volatility, etc. For a large index like the S&P 500, λ might be on the order of 1e-6 to 1e-5 per dollar – meaning, for example, that an extra \$1 billion of buy pressure could move the index by 1–10% if not absorbed [48]. This translates to a dimensionless **impact intensity** κ = S·λ, which might range from 0.1%

to a few percent for the index [48]. In normal conditions, κ is small (the market can absorb flows); in stressed conditions, κ rises (market moves a lot per unit flow).

The gamma and vanna fields interact with liquidity through such a parameter. If dealers' required hedge flow Q (in dollars per unit time or per ΔS) is large relative to market volume, the hedging itself will move prices. We can conceptualize a modified price dynamics equation that incorporates these fields. One such formulation (from an *enhanced hedge pressure* model) is:

$$(1-\kappa_t\Psi_t)\,dS_t = \kappa_t\,w_{CH}\,\Phi_{CH}(t)\,dt + \mu\,dt + \sigma\,dW_t\,,$$

where $\Psi_t = w_\Gamma\,\Phi_\Gamma(t) + w_{Vanna}\,\Phi_{Vanna}(t)$ is a combined measure of gamma and vanna "pressure" at time t, and $w_\Gamma, w_{Vanna}, w_{CH}$ are weighting factors calibrating how much each exposure translates to actual flow [49] [50]. In this equation, $(1-\kappa\Psi)$ acts like a **feedback factor** reducing or amplifying the effective drift/volatility of S. If Ψ_t is positive (dealers long gamma/vanna such that they buffer price moves), then $\kappa\Psi$ subtracts from 1 – making the coefficient on dS_t smaller, i.e. price changes per unit noise are smaller. In the extreme, if $\kappa\Psi \to 1$, the denominator $(1-\kappa\Psi)\to 0$ and the **amplification factor** $A_t = 1/(1-\kappa\Psi_t)$ blows up [51]. This would correspond to a scenario where the market's own hedging feedback is strong enough to create instability (a theoretical tipping point where dealer flows could completely overwhelm liquidity or vice versa). In practice, risk systems monitor this to ensure it stays well below 1. For example, one production hedge engine requires that the saturated pressure $|\kappa\Psi|$ stays below 0.8 with >99% probability [52] – essentially ensuring a safety margin so that feedback loops don't get too close to runaway. On the right-hand side of the equation, the charm term $\kappa\,w_{CH}\,\Phi_{CH}\,dt$ indicates how time-decay flow (charm) introduces a *drift* in returns (positive or negative depending on sign) [49]. This formalism merges market impact with the Greek fields to yield an **"effective" price process** that is *endogenous*: part of the drift/volatility comes from the market's own positioning and hedging activities.

From a **market regime** standpoint, gamma, vanna, and charm fields help explain why volatility tends to cluster in regimes. In long stretches of calm, bull-market conditions (e.g. 2017 or late 2019), we often observe high levels of positive dealer gamma exposure – vol is low, investors sell options for yield, and dealers are long gamma. This creates a virtuous cycle: dealers' long gamma hedging dampens realized volatility, which in turn encourages more option selling (keeping implied vols low), and so on. Indeed, it has been documented that realized volatility stays very muted when aggregate gamma is strongly positive [17], and that realized vol only spikes once that positive gamma supply is removed or flips sign [16]. In these low-vol regimes, the vanna field is usually benign as well – implied vol is low and often falling slowly (declining IV adds supportive flow as discussed), and any spot-vol feedback is mild. Liquidity is abundant (κ is low), so even if there are hedging flows, they have little impact. The net effect is a **stable equilibrium** where both implied and realized volatility are suppressed and option hedging flows reinforce the stability (often, realized vol << implied vol, leading to profits for option sellers a.k.a. volatility risk premia).

On the other hand, in stressed regimes or bear markets, conditions invert. Customers tend to buy protection (long puts) and unwind short calls, making dealers short gamma and short vega. The gamma field goes negative (or at least much smaller), meaning the market loses that stabilizing buffer or even gains a destabilizing feedback. At the same time, liquidity (κ) often deteriorates – e.g. fewer resting orders, wider spreads, or outright withdrawal of liquidity providers – so any hedging flows move price more. The vanna field becomes treacherous: as the market sells off, implied vol jumps (spot-vol correlation flips to highly negative), unleashing vanna flows that *add fuel* to the sell-off (dealers needing to sell more as vol rises). This is why we see phenomena like **autocorrelation of returns and "momentum" during sell-offs** –

studies have found that short-gamma hedging indeed *creates* intra-day momentum when dealers are offside [53] [54]. If the sell-off is severe, a vicious cycle can form: dealers' short gamma means they're selling into declines, which pushes prices lower, which increases vol, which via short vanna means they must sell even more. This positive feedback loop can result in cascading losses and liquidity crises. Historical crashes like 1987's Black Monday have been partially attributed to dynamic hedging strategies (portfolio insurance) which are conceptually similar to widespread short-gamma positions that had to be unwound into a falling market. More recently, events like the 2018 "Volmageddon" (VIX ETN blow-up) or the March 2020 sell-off saw dealer positioning exacerbate moves: in early 2018, many dealers were short volatility through VIX products and calls, and when volatility spiked, they had to buy VIX futures and sell S&P futures, accelerating the turmoil. In March 2020, markets saw implied vol go to extreme highs while dealers were caught short puts, creating enormous negative vanna such that even small bounces were sold into until the positioning cleared. Essentially, the **market regime flips from mean-reverting to trending when the gamma field changes sign**, especially if coincident with a volatility spike that activates vanna flows [55].

Liquidity and volatility also interact with these fields via the **implied volatility surface**. As part of modeling, one must know how implied vol itself changes with underlying price (the "volatility skew"). Dealers' vanna exposures are linked to the skew: e.g. if the market falls, not only does overall IV rise, but the skew means OTM puts gain even more implied vol. A well-known dynamic is that the skew can soften the blow of gamma somewhat – as the market falls and dealers sell stock (short gamma), the rising IV increases their options' values and deltas, which is *partially* hedged by that selling (in a sense, some of the move is absorbed by option revaluation). However, that same skew dynamic means dealers' vega risk (exposure to implied vol) grows – which can prompt them to hedge by **buying volatility (VIX futures or variance swaps)** if they are short vega. This is another feedback: a dealer short a lot of index puts might hedge the vega by buying VIX futures; if the market drops, they will buy more VIX (pushing implied vol up further) which then increases their delta (via vanna) requiring more stock selling – a cross-market feedback loop. Such interactions show that one cannot view gamma, vanna, charm in isolation; they operate within a complex ecosystem of **liquidity provision, volatility pricing, and cross-asset flows**. Market regimes (low-vol calm vs high-vol crisis) essentially correspond to different dominant feedback loops: in calm times, negative feedback loops dominate (stabilizing forces: long gamma hedging, vol selling strategies, etc.), whereas in crises, positive feedback loops dominate (destabilizing forces: short gamma hedging, margin calls, etc.).

An interesting case study is the rapid growth of **0DTE options in 2023-2025**. These ultra-short options have huge gamma and charm (because τ is near zero) but essentially no vanna beyond the trading day. Many feared that 0DTE trading by speculators might cause intraday volatility spikes. However, research shows that market makers in 0DTE are typically *net long gamma on average* (they sell these very short options to intraday traders), which actually *reduces* volatility and promotes mean reversion intraday [20]. Consistent with theory, days with particularly high customer demand for 0DTE calls or puts can lead to episodes where dealers get short gamma intraday (if, say, everyone buys calls in a sudden rally, dealers short those calls will be short gamma for a few hours until they can flatten out). During those moments, we have seen abrupt swings or "intraday momentum" bursts. But because expiration is at the close, any gamma imbalance resets the next day. In effect, 0DTE options have shifted some of the gamma field's influence from multi-day swings to *intraday* patterns. The charm on 0DTEs is extreme (all the delta decays to zero by end of day), which means there is reliably a large hedging flow in the closing hour as those positions roll off. Traders have adapted to this, and market makers dynamically adjust all day, so thus far the market has absorbed 0DTE flows without a major incident. It's a testament to how robust the dealer hedging framework is – and how critical it is for institutions to monitor these exposures in real time.

Implementation in Modern Systems: From Implied Order Books to Hedge Engines

Both sell-side institutions (banks, market-makers) and sophisticated buy-side firms (hedge funds, prop desks) implement the above concepts into real-time systems. Two notable implementations are: (1) **Implied order book models** that map out gamma and vanna fields to predict support/resistance and liquidity pockets, and (2) **Dynamic hedge engines** that integrate these fields with market impact and risk management for automated trading or risk hedging.

On the research/analytics side, the *implied order book* idea has become popular among volatility traders. By analyzing options position data, one can compute Dealer Directional Open Interest (who is long/short each strike) and then calculate how all those options' deltas would change for small moves or vol shifts [56] [57]. Plugging in a range of S (spot) values into Black–Scholes formulas yields the aggregate delta change (per point move) = gamma exposure profile, and plugging in different vol values yields the aggregate delta change per vol point = vanna profile [58] [59]. These profiles essentially act like **supply and demand schedules** for the index. For example, if the SPX is at 4000 and there is a massive positive gamma exposure from 3950 to 4050 strikes, then if SPX drops to 3980, dealers will start buying in size (providing support), and if it rises to 4020, dealers will sell (capping the rally). This can be visualized as a **"gamma ladder"** or heatmap across strikes. Traders use services that publish metrics like **Gamma Exposure (GEX)** at each strike and the **"Gamma flip" level** (where net gamma changes sign). These help anticipate when the market might transition from stable to volatile. Similarly, vanna exposure can tell you how sensitive dealer hedging is to a vol shock – for instance, if a particular expiration has a lot of short put positions for dealers, a spike in IV could trigger a big sell program from dealers (as we saw in 2020). By knowing this, a firm might pre-emptively reduce risk or even speculate on the vol-event by, say, buying protection if they know dealers are dangerously short vanna.

Many institutional systems incorporate these calculations internally. For example, an **"Enhanced Dealer Hedge Pressure Field Engine"** described by one quant group provides a production-ready framework: it continuously ingests options market data (prices, volumes, OI), estimates the **arbitrage-free implied vol surface** (using a parametrization like SSVI to fill gaps and ensure smoothness [60] [61]), computes Greeks for all options (with proper unit handling, e.g. converting charm to per-minute decay) [3] [62], and then aggregates the Φ_Γ, Φ_{Vanna}, Φ_{Charm} fields across the chain [9] [10]. Crucially, it also **determines the sign of these exposures from the perspective of dealer hedging** – which requires knowing whether dealers are net long or short each option. This is done via **trade classification algorithms** (like Lee-Ready tick tests on options trades, combined with open interest changes) to infer who the initiator was and thus whether a trade added to dealer long or short positions [56] [63]. Once the fields are computed, the system doesn't just assume they impact price in a vacuum; it uses **direct market impact measurements** to calibrate their effect. Specifically, it measures Kyle's λ in real time from order flow data [46] [47], so it knows, for example, "right now, \$100M of hedging flow would move the index by X points". With this, the system produces *adjusted predictions*: it can forecast short-term price moves or volatility by combining the pressure fields with the liquidity metric. For instance, it might output an expected 5-minute return given the current charm flow and an amplification factor accounting for gamma/vanna feedback [49] [51]. It even computes an **"amplification factor"** as described earlier, and a **"decay wind"** to quantify charm's influence [64]. All this is delivered via an API so that trading algorithms or risk dashboards can consume it [65] [66].

In practice, what might a trading desk do with such a system? A few applications:
- **Dynamic Delta Hedging and Market Making:** A market-maker can feed these signals into their auto-hedger. For example, if the model indicates that the market is in a short-gamma regime (Ψ is large and negative) and amplification factor is high, the market-maker knows price moves will be exaggerated. They may widen spreads or reduce position size, since any inventory could lead to fast losses. Conversely, if long gamma regime, they can trade more aggressively mean-reverting, knowing moves will be dampened. Some firms also use this to *optimize their own hedging*: rather than hedging every small move (which can be costly), they might tolerate more inventory in a long-gamma environment because the risk of a large move is lower [17]. In short-gamma environments, they might hedge more frequently or even overshoot hedges anticipating momentum.
- **"Implied Order Book" Visualization:** Institutional traders may overlay the gamma field on price charts – showing key strikes as "liquidity levels". This helps portfolio managers decide where to place stops or take profits (e.g. avoiding getting stopped out in a low liquidity air-pocket just beyond a big strike). It's also used in options trading strategies like pinning plays (betting the stock will close near a strike) or in understanding why markets suddenly accelerate after breaching certain levels (often, it's because they moved out of a high-gamma region into a low-gamma one – the proverbial air pocket).
- **Portfolio Hedging and Tail Risk:** Risk managers monitor aggregate gamma/vanna across the book. If they see the firm is net short a lot of gamma heading into a potentially volatile event (earnings, Fed meeting), they might buy back some options or otherwise hedge to avoid being forced to liquidate into a fast market. Some hedge funds explicitly run strategies that **harvest the gamma field** – e.g. buying options when the market is net short gamma (expecting an outsized move to come) and selling options when the market is net long gamma (capturing premium during the calm). These strategies are informed by metrics like total Gamma Exposure (GEX) and Vanna Exposure (VEX). The earlier-cited paper by Boyd et al. even suggests that policymakers could watch net gamma positioning as a financial stability indicator, and hypothetically take actions to alter it (though how is debatable) [19].

From a *systems* perspective, implementing these models requires robust data and risk controls. The "Dealer Hedge Pressure Engine" document emphasizes things like **arbitrage-free vol fitting** (to not mis-estimate Greeks) [60], **real-time updates** (e.g. recalibrating λ every minute, and refreshing regime probabilities hourly) [67], and **temporal data integrity** (no using future data – e.g. yesterday's end-of-day OI – when predicting intraday) [68]. It also highlights **unit consistency and sign conventions**: e.g. ensuring that sums like $S \cdot \Phi_\Gamma$ are scale-invariant (so the model works if the underlying price changes or if measured per index point vs percentage) [69], converting charm to per-minute [6], and using the correct sign for dealer vs customer positions [9] [10] (since a customer long call = dealer short call, which flips the sign of gamma exposure from the dealer's perspective). Many such engines include **safety checks**: for instance, if the vol surface quality is poor or OI data incomplete, confidence in outputs is lowered [70], and perhaps the system reverts to a simpler baseline. Some also incorporate **regime classification** – e.g. outputting the probability that we're in a short-gamma regime vs long-gamma regime based on current Ψ, vega trend, etc [71] [72]. This can be used to trigger different trading strategies or alerts.

Lastly, some large institutions integrate these models into an **automated hedging or "volatility control" strategy**. For example, a dealer with a large options portfolio could automatically delta-hedge more passively or aggressively depending on the pressure fields: if model says high amplification (dangerous feedback), it might break hedges into smaller slices over time (to not shock the market), whereas in benign conditions it might do larger, less frequent adjustments. Buy-side funds that *don't* make markets also pay attention, because dealer flows can create opportunities or risks for their positions. A long-only fund might reduce equity exposure ahead of a known index option expiry that is predicted (via gamma maps) to remove a lot of support from the market. Alternatively, a short-term trader might buy an index future right into the close of expiration Friday if they know a large charm-driven buy flow is coming as dealers unwind short hedges.

In summary, today's markets are increasingly **"options-aware."** The once esoteric Greeks gamma, vanna, charm are now widely monitored not just by quants but by discretionary traders, through reports and dashboards. Firms like SpotGamma and SqueezeMetrics popularized charts of gamma exposure and "vanna/charm flow" commentary, which have even made it into financial media. Meanwhile, on the back end, high-tech hedge engines implement these calculations rigorously, ensuring that trading decisions are informed by the full landscape of time–price Greek fields and not by price action alone.

Limitations and Pitfalls of Field-Based Models

Despite the power of gamma, vanna, and charm fields in explaining and predicting market behavior, there are important limitations and potential pitfalls in relying on these *field-based models*.

1. Model Risk and Assumptions: At their core, these models often assume Black–Scholes-like behavior – continuous prices, continuous hedging, and a certain relationship between underlying moves and volatility (often a static skew or a stochastic vol model). Reality can violate these assumptions. Markets can gap (discontinuous jumps), instantly rendering dynamic hedging models inaccurate. For example, a huge overnight move can bypass the "hedging support" that gamma fields might imply during normal trading. Likewise, implied volatility can change not just as a deterministic function of price and time, but due to swings in sentiment or liquidity that are outside the model. If traders put too much confidence in a gamma map, they might be blindsided by a regime change. The 2010 Flash Crash, for instance, was not predicted by any gamma model – it was a liquidity shock from other causes.

2. Position Transparency: These models are only as good as the position data underlying them. For index options, one can estimate dealer vs customer positioning from trade and OI data (as done with DDOI – Dealer Directional OI – in the implied order book approach) [56] [63] . But there's noise and potential error in that process. If the sign of positions is mis-inferred, the computed fields could have the wrong sign (thinking dealers long gamma when in fact they are short, etc.). Furthermore, not all players hedge in the same way: a hedge fund might buy calls from a market maker and that market maker delta-hedges (classic scenario). But another fund could write calls and *not* delta-hedge (taking outright directional risk). The models usually assume **dealers hedge, customers do not**. If customers also dynamically hedge or if dealers choose to run a directional book (which sometimes they do under risk limits), the flows might differ from model predictions. The field approach tends to lump all "dealer" behavior together as if it's uniform; in practice, some dealers may have better or worse execution, timing, or may *pre-hedge* in anticipation of flows. These nuances can cause deviations between predicted and actual market impact.

3. Liquidity Regimes: The parameter κ (Kyle's lambda) can change rapidly, and in extreme cases market liquidity can dry up completely. Field models typically linearize the impact – assuming $dS = \lambda\, dQ$ in some fashion. In a crisis, it might become highly non-linear (e.g. each additional sell order has a bigger impact than the last as book depth vanishes). This means that negative feedback from gamma could suddenly fail. For instance, dealers might be happily providing liquidity on the way down, until they hit position or risk limits – at which point they stop hedging (or even start liquidating), and the remaining gamma support evaporates. A known pitfall is **assuming continuous hedging is always possible**: if trading halts or if moves are so fast that hedging lags, the theoretical stabilizing effect of long gamma is less than expected.

The 1987 crash is a case where the intended dynamic hedging (portfolio insurance, analogous to short put hedging) couldn't keep up with the market – once it fell too far too fast, the hedges all came in at once and there were no buyers.

4. Vanna/Charm Complexity: Vanna in particular is complex because it's entangled with the volatility surface. Most models treat vanna as a linear sensitivity for small vol changes. But in a true vol shock, the entire skew can change shape, and second-order Greeks like **vomma** (sensitivity of vega to vol) or **dvega/dtime** can matter. A purely gamma-vanna-charm field model might misjudge flows if, say, implied vol doubles – a lot of assumptions (like constant vanna or correlation) break in that scenario. Charm, while predictable, can be overwhelmed by other flows easily. It's a slow burn, and if a big price shock hits, the charm-induced drift is negligible by comparison. So one pitfall is **over-emphasizing charm**: indeed some analyses just ignore it as noise [41]. The truth is charm is usually second-order, except possibly right *on* expiration, and even then other factors (like traders rolling positions or closing them) can dominate over pure time-decay hedging.

5. New Option Flows: The fields often treat the existing open interest as static and project its effects. In reality, open interest itself changes as new trades happen. For example, the models might predict that after a big expiration, volatility will increase (since gamma support drops). This is often true *ceteris paribus*. But the day after expiration, new options could be traded that add fresh gamma exposure (e.g. traders might initiate new positions in the next expiry, reintroducing gamma into the market). Or, if the market does start moving post-expiry, the surge in implied vol might lead to lots of new option demand (further altering positions). In other words, the **assumption of fixed positions can break down in a feedback way**: a big move will itself spur option trading that changes the fields. This makes the system highly dynamic and sometimes reflexive – e.g. a rapid drop might cause call sellers to emerge (adding positive gamma) or put buyers to pile in (adding more negative gamma); the model would need to update rapidly to reflect that.

6. Cross-market and Exogenous Factors: Option-derived pressure is just one source of market moves. There are also fundamental flows (earnings, macro data, buybacks, mutual fund flows, etc.) and other systematic strategies (like volatility targeting funds, CTAs, risk parity) that can generate buying or selling independent of options. Sometimes these other flows dominate, and an option field model might falsely attribute a move to gamma when it was really caused by, say, a central bank surprise. Conversely, a move predicted by the model may not occur if an exogenous big buyer or seller steps in against it. For example, the model might say "dealers will have to sell $5B if market drops 2%" – normally that suggests a possible cascade. But if a pension fund or the Fed or some large actor decides to buy into that drop, they can absorb it and break the loop. Thus, **field models should be used as one input, not a sole determinant**.

7. Implementation Challenges: On the practical side, calculating these fields in real time with accuracy is non-trivial. Data issues like stale or erroneous open interest, corporate actions (e.g. an index rebalancing or special dividend changing option parameters), or illiquid strikes can all introduce errors. There is also a pitfall in **double-counting or overlap** – gamma and vanna effects can coincide and one must ensure not to over-estimate combined impact (this is why the advanced model uses weights w_Γ, w_{Vanna} and even an uncertainty overlay to calibrate how much of observed moves are attributable to each factor [72]). If not calibrated, a naive model might sum gamma and vanna impacts and overshoot actual price effects.

8. Rare events and Tail Risk: Field-based models are often built on historical patterns (e.g. regression of flows to returns). They work until they don't – usually failing in tail events that are by definition not well-represented in history. A known phenomenon is **"gamma trap"** or **"negative gamma spiral"**: when dealers are caught massively short gamma and the market starts tanking, the usual relationships can invert (implied vol can shoot up so much that eventually dealers stop selling stock and start *buying volatility at any price*, causing option prices to explode and liquidity to vanish). These kinds of tail outcomes are hard to incorporate. Often the models will signal increasing risk (e.g. amplification factor rising), but *quantitatively* they might still understate how fast things can unravel once certain thresholds are passed.

In conclusion, while gamma, vanna, and charm field analyses have proven extremely useful – offering insight into the "invisible hand" of dealer hedging – one must use them with a good understanding of their limits. They excel at explaining the **aggregate, predictable component** of flows, but they cannot foresee everything. The market is a superposition of many forces, and the options-derived fields are just one layer (albeit a significant one in modern markets). Prudent practitioners combine field models with other analyses, maintain healthy skepticism, and stress test scenarios (e.g. "what if implied vol jumps 10 points – how would our model behave?"). As with any model, human judgment and vigilance are required, especially as markets evolve (for instance, the rise of 0DTE options required modelers to adjust how they account for charm and intraday gamma).

Conclusion

Viewing options Greeks as dynamic fields over a time–price landscape provides a powerful conceptual and quantitative framework for understanding market price movements and volatility. **Gamma** acts as a stabilizing or destabilizing field depending on its sign: a high positive gamma field means dealers supply liquidity and dampen volatility, whereas a negative gamma field can lead to self-reinforcing volatility and trending moves [18] [16]. **Vanna** connects volatility and price in a feedback loop: a benign vanna environment can augment stability (vol declines leading to dealer buying), but a highly negative vanna field in a crisis can dramatically amplify sell-offs as volatility spikes force further dealer selling [35] [36]. **Charm** provides a more subtle temporal pressure, often manifesting as a gentle drift due to time decay that becomes noticeable around option expirations [40]. Together, these fields create a mosaic of "pressure" on the market – one that traders and models increasingly monitor to gauge where the market might resist or accelerate.

The interplay of these Greeks with **liquidity and market regimes** explains many empirical phenomena, from the calm of a pinned market to the violence of a volatility cascade. High positive gamma positioning corresponds to tight intraday ranges and rapid mean-reversion (as evidenced by both observational studies and 0DTE data) [17] [20], whereas negative gamma positioning is linked to volatile swings and momentum [18] [20]. Vanna and charm flows further modulate these patterns, adding complexity especially in how markets respond to changing volatility or the passage of time.

Modern institutional systems have operationalized these insights. Dealers and hedge funds build real-time "**implied order books**" from gamma and vanna distributions [21], and **dynamic hedging engines** that incorporate measured market impact ensure that theoretical pressures translate to practical trading signals [73] [49]. These systems allow market participants to anticipate when dealer hedging will provide support or exacerbate a move, and to adjust their strategies accordingly – for example, scaling back risk when amplification factors are high, or harvesting risk premium when the fields indicate a stable regime.

Yet, we must remember that these models are abstractions. They capture an important, sometimes dominant, piece of the market's mechanics, but not the whole thing. They assume a certain structure to behavior (dynamic hedging) that generally holds but can break in extreme conditions or when other players intervene. As the adage goes, “all models are wrong, but some are useful.” The gamma, vanna, and charm field paradigm has proven extremely useful – turning esoteric Greeks into actionable market maps – but it must be applied with care to its assumptions and in concert with a broader market understanding.

In summary, gamma, vanna, and charm as dynamic fields offer a conceptual lens and practical toolkit for navigating the ever-evolving landscape of modern markets. They illuminate how the “derivatives forces” shape price action across time and price, highlighting feedback loops that can either stabilize or destabilize the system. By studying these fields, both theorists and practitioners gain a deeper appreciation for the connected nature of options and equities, and why the pricing of risk (implied volatility) cannot be divorced from the actual risk (realized volatility) in the market. Armed with this knowledge – and mindful of its limits – traders and risk managers can better anticipate liquidity crunches, volatility spikes, or, conversely, those eerie calm periods where the market seems to gravitate toward a point. The time–price landscape, crisscrossed by Greek-derived fields, is a rich terrain that continues to be mapped and explored, offering insights into the hidden order behind market disorder.

---

1 21 22 23 24 25 26 30 31 32 33 34 35 36 37 38 41 56 57 58 59 63 The_Implied_Order_Book.pdf
file://file-A59WsbyAM8UeVnMd1G63H4

2 0DTE Options Charm Decay: Expiration Day Trading Guide
https://www.aiflowtrader.com/blog/charm-0dte-dynamics

3 4 5 6 7 8 9 10 11 12 39 43 44 45 46 47 48 49 50 51 52 60 61 62 64 65 66 67 68 69 70 71 72 73 Options_Pressure_Field.pdf
file://file-9HSki1XRjLooNRrTmvtDyR

13 14 15 27 28 29 40 42 Dealer Hedging Mechanics - Menthor Q
https://menthorq.com/guide/dealer-hedging-mechanics/

16 17 Nomura’s McElligott Explains The ‘Perfect Virtuous Feedback Loop’ As Gamma Gravity Pins Equities "Intraday movement is squelched". | SpotGamma™
https://spotgamma.com/nomuras-mcelligott-explains-the-perfect-virtuous-feedback-loop-as-gamma-gravity-pins-equities-intraday-movement-is-squelched/

18 19 55 Gamma Positioning and Market Quality by Boyd Buis, Mary Pieterse-Bloem, Willem F. C. Verschoor, Remco C. J. Zwinkels :: SSRN
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4109301

20 54 0DTEs: Trading, Gamma Risk and Volatility Propagation by Chukwuma Dim, Bjorn Eraker, Grigory Vilkov :: SSRN
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4692190

53 Hedging demand and market intraday momentum | Request PDF
https://www.researchgate.net/publication/351315710_Hedging_demand_and_market_intraday_momentum